

Title to be revisited

Nachtrab Sean

SINF - Computer science
Medical informatics

Supervisors:

Prof. Greet Kerckhofs
Prof. Sébastien Jodogne

Readers:

Isabelle Gennart
Juliette Vanderhaeghen

© This template was adapted by the author from the *Typst KUL template*

Without written permission of the supervisors and the author it is forbidden to reproduce or adapt in any form or by any means any part of this publication.

A written permission of the supervisors is also required to use the methods, products, schematics and programmes described in this work for industrial or commercial use, and for submitting this publication in scientific contests.

Preface

This thesis was launched at the initiative of the writer wishing, in the context of a degree in medical informatics, to work on a problem at the crossroads of two fields, and as a result has two supervisors. It extended work done during a previous thesis, and made use of data collected in suboptimal conditions.

The developped plugin is available at [TODO](#) upload to UCL git, with the written report code available on [github](#)

The Generative AI tool Piccolo was used during the writing of this thesis for feedback and writing modifications for certain sections, although no sections were fully written using AI, for a total use of xxx credits as detailed further in Appendix [A.3](#). Claude Code was used for support during certain sections of software development, mainly pertaining to the initial setup and 3DSlicer point placement interaction.

Nachtrab Sean

Contents

Preface	i
Abstract	iv
List of Figures	v
List of Abbreviations and Symbols	vi
1 Introduction	1
1.1 Biological motivation	1
1.2 The role of computer science	2
1.3 Computer science as a tool	3
2 State of the art	5
2.1 Software in Research	5
2.2 Tissue imaging	8
2.3 Vascular Segmentation	10
3 Problem statement	17
3.1 Problem resolution approach	17
4 Software tool development	19
4.1 3D Slicer Plugin	19
4.2 Segmentation of tumor vascularization	21
4.3 Results	24
5 Contending with limitations	27
5.1 Micro-CT Data management challenges	27
5.2 Performance challenges	27
5.3 Ablation study	30
5.4 Conclusion of the ablation study	30
6 Micro vasculature extraction	33
6.1 Final vessel extraction flow	33
6.2 Data annotation for evaluation	33
6.3 Performance results	33
7 Conclusion and future perspectives	35
A Appendices	39
A.1 Environmental and CO2 impact of the thesis	39
A.2 Instructions for scientific rigor in data handling	40

A.3 Use of LLMs and AIs	40
A.4 Structuring a Masters thesis	40
A.5 Masters thesis writing across group and field boundaries	40
B Extra figures	43
Bibliography	45

Abstract

TODO

Driven by the link between tissue structure and function, medical imaging has evolved from its humble beginnings of 2D microscopy towards higher and higher resolution 3D imaging. Contrast-enhanced micro-tomography, one of these high resolution methods, was used with the purpose of separating the micro-vasculature of tumors,

The process of data analysis, its associated tools and methods, has become central to the scientific process.

Advanced imaging methods such as contrast-enhanced micro tomography aim to separate tissue by use of contrast enhancing agents, but noise in imaging and the multiscale, dispersed nature of vascularization result in the need for more powerful software methods to enable the comparison of samples across various treatment profiles, and achieve the levels of accuracy required for drug research. In this document the clinical pipeline is examined from end to end, a more powerful extraction method for vascular structures in challenging use cases is developed, and the groundwork laid for principled, scientifically rigorous work to be carried out.

List of Figures

3.1 Preliminary synthesized diagram for the annotation of vascularization and sharing of results	18
4.1 Left: , Right: QT user interface designer	20
4.2 Structure for module implementation from [1].	21
4.3 View of the seed annotation and vessel size definition panes. The user may press add, click on the locations in any of the right hand panes, as well as import previous annotations or export the current points.	22
4.4 Illustration of the shortcoming of threshold based segmentation, with the “shell” being included	23
4.5 Visualization of the vesselness probability map, as defined by the Frangi feature. Higher valued pixels are assigned a higher vesselness probability.	24
4.6 Visualization of the segmentation following the multi stage pipeline during development	25
5.1 TODO: This is a placeholder. TODO: Add horizontal lines for the RAM of computers. Visualization of the raw dataset sizes, obtained by multiplying width, height and depth by 8 bits per pixel	28
5.2 1. CA-LL-L2 RAM utilization during segmentation: baseline after loading dataset, bellow 24GB utilization, during processing 100% RAM and swap are used, 2. Run 1, CB-LURU-R showing full RAM and SWAP utilization	29

List of Abbreviations and Symbols

Abbreviations

CT	Computed tomography, method for 3D imaging
Micro-CT	High resolution CT
CECT	Contrast-Enhanced Micro-CT
CESA	Contrast-Enhancing Staining agent
CCA	Casting Contrast Agent
GT	Ground Truth

Chapter 1

Introduction

In the field of biology the cutting edge in analysis of tissues has evolved from 2D histology, carried out by humans looking through microscopes, to a wide variety of 3D techniques across different wavelengths, ranging from sound and infrared to magnetic resonance and x-ray imaging. The pursuit of higher resolution and fidelity reconstruction is driven by the importance of tissue structure in function, particularly for vascular networks where it plays a critical role. While 2D histology is established as the gold standard for tissue analysis, its limitations in capturing three dimensional structure have driven the adoption of high-resolution 3D imaging techniques, such as micro-computed tomography (Micro-CT) and its contrast-enhanced variants aiming to improve tissue separation. These methods generate orders of magnitude more data than traditional microscopy, enabling detailed analysis of complex structures like tumor vasculature but posing new challenges in extraction and processing. Small vessels are characterized by low contrast, noise, and contrast-enhancement agent gradients, not able to be accurately reconstructed using conventional methods. Current approaches also suffer from poor replicability and accessibility due to the widespread use of proprietary software. This gap between computational advances and biological workflows underscores the need for open-source, user-friendly tools integrating principled extraction methods.

To address this a pipeline for micro vasculature extraction from CECT data is developed in the form of a plugin for 3D Slicer, an open-source tool widely used in the medical imaging field. Scientific rigor is enforced by removing subjective user adjustment of hyperparameters, robustness and transferrability to new data is improved by using algorithms with adjustable hyperparameters, and adoption is encouraged by optimizing usability, to bridge the divide between computer vision research and biological applications.

1.1 Biological motivation

Biological tissues are spatially complex, diverse in composition and organization of their constituents at the microscale. Tissue microstructure is closely linked to its function: soft tissue analysis is foundational to modern medicine, being used everywhere from clinical environments to laboratories. Tissue samples are collected and imaged in order to gain an understanding of their structure and composition, called histology or histopathology in the case of diseased tissue. Classically, this sampling and subsequent analysis has been done *ex-vivo* in 2D using a microscope, often manually analyzed, setting the gold standard for

histology [2]. However classical 2D histology presents the distinct disadvantage of being unable to obtain information at the same granularity in all three axis: slices may be stacked to obtain a 3D result, but resolution in the third axis is limited to section thickness, with the stacked planes separated by a dead space where no information is acquired [3], and the cutting process induces changes in tissue structure [4].

For the purpose of evaluating the use of the antiangiogenic drug Pazopanib, a tyrosine kinase inhibitor which blocks the creation of new blood vessels as a means for cancer treatment, accurate reconstruction of tumor vasculature is required. Vascularization is by nature 3D and covers entire tissues, and simple volume estimates do not capture important parameters such as tortuosity, connectivity/branching and cross section profile. This 3D nature, coupled with the limitations of 2D methods, motivates the use of 3D imaging: for the purpose of tumor evaluation, contrast-enhanced Micro-CT scans were collected with an agent to increase the contrast of blood vessels.

1.2 The role of computer science

Microscopy as a technique and method for tissue analysis predates the field of computer science. As noted, 2D microscopy has downsides and use cases where tissue structure is of relevance have pushed the use of more data intensive and computerized imaging techniques, namely 3D methods. To collect 3D data, the simplest is to extend 2D techniques by stacking slices [5]. This method limits the resolution across the slicing axis, and does not alleviate the tissue deformation challenges, although is doable without extensive computerization. Non invasive 3D imaging methods rely on wavelengths able to further penetrate the tissue to be imaged: visible light can be used for shallow 3D imaging of soft tissues, with longer wavelengths such as x-rays able to penetrate completely through soft tissue as well as cartilage and bone, requiring a computerized reconstruction pipeline to generate final usable data. Although the machines and methods exist to collect high resolution 3D data using various wavelengths and associated techniques, such as Micro-MRI and Micro-CT, the analysis of high resolution 3D data is not as straightforward as with 2D histological slices: there are many more slices, often thousands, requiring exponentially more time for manual review and annotation.

There is also the issue of data diversity stemming from the increase in dimensionality and available imaging methods: even for one given imaging type such as Micro-CT, there is high scan to scan variance between different tissues, a wide array of machines and machine settings, and different contrast or staining agents that can be used at different concentrations and with different methodologies for getting them into the tissue of interest. This huge diversity naturally drives a huge diversity in data with as consequence a fragmented and diverse data processing landscape.

Our imaging method of choice for analyzing blood vessels, Micro-CT, presents low contrast between soft tissue: techniques exist to increase contrast and improve visibility such as cryo-CT [6] or the utilization of contrast enhancing staining agents (CESAs) in Contrast-Enhanced CT called CECT, as used for the data in this thesis. Although essential to increase the contrast between tissue types, these techniques still do not enable sufficient differentiation

for small vascular networks such as those found in the murine tumors collected for the evaluation of Pazopanib to allow threshold based classification as is commonly used for bone or for samples with high contrast agent presence [3], this was shown during the analysis phase of [7], the source of the data used here, where threshold based segmentation was used and resulted in poor performance.

1.3 Computer science as a tool

When faced with complex data the paths for researchers are twofold: either manual annotation of the data, or utilization of algorithms to aid in the separation of their structures of interest. Manual annotation suffers from high subjectivity, strong variance between annotators [8] and requires time and expert annotators. The inter annotator variance can be explained in multiple fashions: expertise difference, tool type and experience, available time, subjectivity and more. Additionally, annotation natively in 3D is challenging, meaning it is difficult to annotate while properly considering all three axis simultaneously, hindering the annotation of small 3D structures like vasculature. This annotation variance, combined with the time required to manually segment, means that the small vascular networks in the provided data have yet to be manually segmented.

The variance in manual annotations, as well as time taken, motivate the need for automated or algorithm based annotation techniques. Users for whom the purpose of analyzing data is not to contribute or work on the process of analysis are pulled towards user friendly and simple methods, where data analysis is a tool to enable information extraction. These algorithmic annotation techniques can themselves again be split into two categories: (1) simple, generalist methods such as thresholding, easily understandable and widely available, however having poor performance on data that cannot easily be split purely based on grey values and (2) a wide spectrum of techniques that carry a richer prior or set of priors about the target structure, ranging from Otsu to machine learning. For vessel analysis, simple methods are easy to apply but result in discontinuities in blood vessels and don't generalize across samples, as well as require the selection of a fixed threshold per sample, which was discovered during user interviews to be carried out based on the users prior experience, not a principled, repeatable manner.

The uptake and adoption of advanced techniques remains limited outside of the walls of computer science labs due to the lack of readily available user friendly solutions, where the barrier to entry or friction associated with using a new tool or method is too high due to a wide array of issues: proprietary paid tools are used, research software is often unmaintained or poorly documented with algorithms challenging to run on different machines due to a large library of pre-requisite software, software versioning, or different OS. For modern deep learning based algorithms, hardware is often a limiting factor due to requiring dedicated GPUs [9], the requirement of particular data formats, or poor performance requiring re-training.

In this thesis work will be carried out to create an open source, user friendly, data driven and robust small blood vessel extraction method for obtaining segmented blood vessels, with

the goal of enabling downstream extraction of clinically relevant vasculature characteristics. This goal is defined by the usecase of a dataset of Contrast-Enhanced Micro-CT data acquired to examine the use of Pazopanib. The pipeline is created as an easily installable plugin for 3DSlicer, an open-source and free base software, that ingests and outputs data in portable data formats familiar to researchers, to ensure re-usability and compatibility with existing downstream software. The plugin implements proven algorithms with user adjustable hyperparameters to enable easy re-use without requiring expertise in computer science. Implemented methods are compared with existing pipelines for extracting fine blood vessels, and tested for robustness by examining performance on data previously considered non segmentable due to non uniform contrast agent staining.

Chapter 2

State of the art

2.1 Software in Research

To work on 3D CT scans researchers at the UCLouvain faculty IMMC (Institute of Mechanics, Materials, and Civil Engineering) use a variety of software explicitly related to the core task of working with 3D volumes: Avizo, CTAn and Dragonfly3D (amongst others), but software is also used implicitly at earlier steps: the OS of the computers, software in the scanner, etc. These pieces of software rarely receive the spotlight in the process of research, but their use places barriers that undermine the scientific principles of open, reproducible, and replicable work by locking processes behind licenses and fees reaching into the tens of thousands of euros per year.

2.1.1 Software Licenses

Software licenses are closely linked with the monetary and scientific costs of use: they govern the terms under which a piece of software may be used, modified, and redistributed and broadly fall into two categories: proprietary (closed source) and open source. Proprietary software restricts access to its source code and is generally distributed under a paid license, although exceptions exist in the form of freeware: proprietary software made available at no cost, as with Dragonfly3D's FreED license. Open source software makes its source code publicly available and is (in the majority of cases) as a result free of charge. Licenses enable the owner of the original software to control the distribution and use of extensions and modifications: Dragonfly3D does not explicitly allow sharing extensions¹, unlike 3DSlicer

Open source software matters in scientific research: First, being free removes a significant barrier to entry, with any researcher, regardless of institutional resources being able to access, run, and reproduce a pipeline built on open source tools. This is a prerequisite for replicability. Secondly open source software is extensible: prior work can be built upon by accessing, modifying and learning from its source code and development. This capacity for incremental improvement mirrors the broader scientific process itself, where each work builds upon prior results. In the medical field, caveats exist to protect from clinical use

¹“[The user] shall not distribute or transfer the Software or Improvements [...], without prior written permission [from Dragonfly3D]” [\[10\]](#)

such as 3D Slicer, published under a modified BSD license. Software such as Orthanc [11] place importance on enforcing sharing of modifications, and is published under a copyleft license, GPLv3.

With this in mind, all code written in the implementation of this thesis, as well as the code to generate this document and the notes taken in the process of redaction, are made open source under the MIT license (for code) and labeled under the creative commons as CC0 - No Rights Reserved (for creative works - the writing).

2.1.2 Software for 3D analysis

The outputs after reconstruction of a CT scan are in the form of a stack of slice-by-slice files that can then be assembled into a 3D volume. Software relevant to this thesis can be divided along the same axis introduced in the previous section: commercial/closed source and open source/free. On the commercial side, the main options include (Bold indicating those actively used in the UCLouvain IMMC) **Dragonfly3D**, **Avizo**, **CTAn**, Amira, Analyze and CTVol (by Bruker), Aphelion, Image Pro, Imaris, MeVisLab, Mimics, ScanIP, Octopus, VG Studio, Zeiss Inspect. On the open source and freeware side, the principal options are 3D Slicer, Chimera, Blender, Dream3D, Drishti, ImageJ, FIJI, IMOD, MeshLab, OsiriX, ParaView, SimVascular, VesselKnife and VisIt. Additionally, code first specialised and focused tools, made in a research context exist such as SPROUT and nnU-Net.

For the purposes of this thesis, software was required to meet the following requirements: **(1)** Import a 3D scan from individual 2D slices in standard formats, **(2)** Export data to non-proprietary formats, and **(3)** Allow coded plugins/code extensions.

With the goal of developing software for practical real world use, the following aspects were also weighed:

1. Availability of dedicated support forums & tutorial videos
2. Active development
3. Prior use by the lab researchers, and by researchers of the field more broadly.
4. Ease of installation and barrier to entry

From this it is clear that standalone approaches requiring coding such as DeepVesselNet [12] and SPROUT [13] are to be discarded as candidates. The software in use at the laboratory for blood vessel segmentation (Avizo, Dragonfly3D) were tested and compared to open source alternatives 3D Slicer and ImageJ/FIJI. ImageJ and its distribution FIJI have a rich history in biological image analysis, having also been used in the past in the lab and on CT data. An extensive plugin ecosystem relevant for this work is available, such as the Frangi vesselness [14] algorithm, but native support for vascular network extraction in 3D is limited.

3D Slicer [15] emerged as the most suitable tool, due to its wide adoption in the medical field, existing plugin ecosystem with vascular specific tools such as The Vascular Modeling Toolkit [16] and R-Vessel-X [17]. Beyond satisfying the core technical criteria of importing 3D scans as stacks of 2D slices, as well as import and export of open formats such as DICOM, 3D slicer also has extensive instructions and prior work in manual and semi-automatic segmentation. The Extension Manager is particularly relevant to this work, enabling non

technical users to install the plugin without needing to code and free of charge, as opposed to the paid extensions for vascular extraction available for tools like Avizo and Dragonfly3D.

Problem 1.

“In an extensive and diverse software landscape, our software must fit into an existing well documented segmentation tool: 3DSlicer, be easy to use: click only installation and offer compelling performance.”

2.2 Tissue imaging

The process of imaging is a critical step in obtaining an understanding of a tissue in both clinical and research contexts. 2D histology is considered the gold standard having a long history [2] and producing high resolution images down to 5 μm for light based methods. The diversity in imaging techniques for 2D histology has rapidly increased the resolution, with electron microscopy able to reach in the order of sub nanometer scale [18]. Progress has also increased the dimensionality of imaging, with computation enabling rapid progress in 3D methods using techniques like tomography, where multiple images are taken non-destructively of a target from different angles and assembled, used in-vivo and ex-vivo methods such as computed tomography (CT) and magnetic resonance (MRI).

For tasks such as measuring the impact of a drug on vascularization, quantification is particularly relevant: Understanding the three-dimensional structure of biological tissue is a prerequisite for the analysis of structure-altering drugs such as Pazopanib, where changes in tumor vascularization are inherently spatial and quantifiable across multiple parameters. For such tasks, 2D histology is insufficient: Physical sectioning of the sample is destructive, is not orientation-agnostic, and introduces deformation artifacts that are difficult to compensate for even with embedding techniques [19], additionally samples may undergo structural changes over time during preparation: they dry out, and certain elements oxidize, although techniques exist to mitigate this [3].

To achieve 3D imaging from 2D histology, 2D slices are stacked across a virtual axis [4]. The resolution achievable when stacking 2D slices to reconstruct a virtual third dimension is limited by slice thickness and inter-slice registration errors. Certain optical techniques making use of slices allow partial recovery of 3D information from 2D acquisitions, such as confocal microscopy, light sheet microscopy, and optical coherence tomography [6], they are limited in sample penetration depth and volume. When the target structure to be imaged and understood does not follow a single preferred axis, as is fundamentally the case for vascular networks and especially those of tumors, the limitations of slice-based histology make proper reliable quantification of blood vessel parameters impossible, and the relevance of 3D methods like MRI and CT clear.

2.2.1 3D imaging techniques for histology

Non-destructive 3D volumetric imaging methods overcome many of the limitations of 2D slice based histology by collecting data uniformly across dimensions, enabling virtual slicing across any plane without requiring physically sectioning the sample. Techniques for 3D imaging range from the aforementioned microscopy techniques using visible or near infrared light, to techniques like MRI that utilize magnetic fields or X-Ray imaging in the form of X-ray microfocus computed tomography (CT). These techniques allow the collection of qualitative and quantitative 3D microstructural data of tissues and their constituents: analysis is not restricted to a single orientation and does not require sample destruction. They also have associated high resolution variants allowing micrometer level imaging: Micro-MRI has an inherently high-contrast for soft-tissue and is able to reach 20 μm [20]; Micro-CT is commonly used at 6 μm , as in the data used in this document and resolutions down to 1 μm

are possible, but the technique suffers from low contrast when imaging soft tissue, when extra techniques are not used [2]. These techniques are the most relevant for ex-vivo tissue histology at the scale of small animal models, due to their small maximum capture volume.

2.2.2 Contrast-enhanced micro-CT

As mentioned, standard Micro-CT suffers from low contrast between soft-tissue making the visualization of blood vessels difficult. To improve contrast two common approaches exist: changes to imaging methodology, and modifications to the sample. Imaging techniques such as phase-contrast Micro-CT (PC-CT) utilize the refractive properties of the X-rays rather than its absorption alone, enhance soft tissue edge detection at the cost of higher complexity and long acquisition times. Tissue modifying techniques that aim to increase contrast include the use of various Contrast Enhancing agents (CESAs): Casting contrast agents (CCA) are perfused into vasculature but pressure must be controlled carefully [21] and is not applicable to our usecase due to the size of vessels to be observed. Diffusion contrast-enhancing agents reach the structure of interest and bind to it based on diffusion through the structure to be imaged, termed CECT. For the tumor vascularization studied in this thesis, ex-vivo tissue binding CESAs are used, namely Hafnium Wells-Dawson Polyoxometalates (Hf-WD POM) due to its nondestructive nature and low tissue shrinkage during incubation after excision. Techniques exist that can be used in combination with CESAs: contrast can be increased using freezing to increase contrast Cryogenic contrast-enhanced microCT (cryo-CECT) which preserves tissue microstructure with reduced deformation [6].

Tumors also present a particularity in that they are frequently partially or entirely devoid of residual hemoglobin, preventing or reducing the action of CESAs and thereby reducing contrast and introducing discontinuities. Diffusion based contrast-enhancing agents also bring an additional disadvantage as seen in the data provided for this thesis: the diffusion of CE agents throughout the tissue happens from the outside in, resulting in a gradient of the amount of agent and as a result a gradient in contrast, as opposed to perfusion CESAs that follow the blood vessels.

Problem 2. *The wide diversity of available modalities for 3D imaging and with their subtypes, the different machines and their acquisition parameters, variability in samples and their preparation, staining agents and methods, as well as large diversity of tissue types come together to present a challenge in creating a method that is re-usable, even within the same lab. Any method used for the segmentation of small blood vessels should be robust to the gradients caused by diffusion CECT, and able to handle variable contrast levels.*

2.3 Vascular Segmentation

The purpose of imaging a tissue is to generate an image containing sufficient data to enable useful information to be extracted. This goal can be seen in the progress of imaging techniques: the improvements in the *acquisition process* laid out in the previous section, to make the data of interest more salient (using agents to enhance the contrast) or to make the data more granular (increasing resolution). The *information extraction* process is explored here and has also seen improvements over time: before widespread computerization, analysis used to be carried out by humans, with rapid iteration since then and evolution in the paradigms used for information extraction.

Vasculature extraction from 3D tomographic data is a longstanding challenge in computer science, extending before the current era of deep learning as seen in [22] reviewing vessel lumen segmentation techniques, although our 3D data carries a unique set of challenges: the vessels of interest are highly diverse ranging in the supplied data from singular voxels to tens of voxels across or 6 to 100+ μm across, corroborated by other researchers doing research using similar methods observing vessels from 10 μm to hundreds of μm [23]. The vasculature is diverse in shape and branching structure, with variable grey levels and most critically disconnections of variable size. They however present a distinctive and consistent set of geometric priors being tubular, connected, branching structures and contain blood which can provide a different imaging profile to the surrounding tissue alongside the CESA in use.

2.3.1 Classical segmentation

Grey value based

In its simplest form, intensity-based thresholding methods such as Otsu's method [24] operate by iteratively finding an optimal intensity threshold that separates foreground from background by minimizing intra-class variance. Such methods are computationally inexpensive and interpretable, but are sensitive to noise, imaging artifacts, and intensity inhomogeneities, as they optimize for a numerical goal. Gradient based methods constitute a natural first approach to CECT, as the goal of utilizing a contrast enhancing agent is to increase the salience of the structures of interest, algorithms that reason globally over an entire image can fail due to diffusion gradients as mentioned previously, as well as noise. Region-growing [25] implementations starting from given point(s) such as marching squares, and watershed algorithms extend this idea by incorporating spatial connectivity, iteratively expanding labeled regions from seed points according to local intensity gradients, making them more robust to global intensity variation, however remaining susceptible to over- or under-segmentation in complex structures. These simple techniques are also not context aware: they can create discontinuities due to not integrating priors of the structure they are segmenting, such as the connectivity of blood vessels.

Geometrically motivated methods

Beyond grey value based methods, algorithms integrating local geometric priors exist such as Gabor filters or Hessian-based filters where the local second-order structure is analyzed

to detect tubular shapes, characterized for vessel detection by Frangi’s multiscale vessel enhancement filter [26], improved in [27]. These second order methods utilize the eigenvalues of the Hessian matrix of local image intensities at multiple Gaussian scales to produce a vesselness score for each voxel, responding to tubular structures while suppressing circular and planar structures. The Frangi vesselness algorithm has been particularly influential in vessel segmentation tasks, as it explicitly models the prior of tubular geometry inherent in blood vessels. These methods offer greater robustness than simple thresholding but depend heavily on the user fine tuning algorithm hyper parameters to optimize performance for a given domain or usecase, and as a result of this feature engineering, can fail to generalize across imaging conditions [14].

Additionally the filter’s response degrades at vessel bifurcations, at the endpoints of vessels, and in regions of low contrast. In [27] extensions to the vesselness formulation are proposed that improve responses at junctions and at low contrast boundaries to increase robustness of the filter.

Extraction can be reformulated to as a global optimization rather than a local filter response to palliate these issues, such as minimal path methods which formulate vessel centerline extraction as an optimization problem. The path of least cost is taken between user-defined endpoints, with cost derived from vesselness or image intensity. In [28] the geometry of vessels is incorporated, adding radius as an additional dimension of the path space. This class of methods is particularly well-suited to expert-in-the-loop approaches where small amounts of data need to be annotated: a user can place seed points that guide the extraction of a complete vascular tree, requiring no training data, and enabling iterative improvement. This method can be seen used for large arteries and airways in [16]. This strength is also a weakness: the low scalability means that for a densely vascularized tissue, or a tissue with non uniform characteristics, many seed points can be required which is impractical. Automatic seed point placement is a potential solution, but remains an open problem.

2.3.2 Data driven learning methods

Data driven approaches represented a major revolution in data analysis over classical methods: they enabled the data characteristics themselves to drive the configuration of the hyperparameters of the algorithm. So called supervised machine learning take in labeled example data and output a result, such as in k-nn a class output based on neighbors, or in tree decision algorithms [29]. The paradigm of deep learning exploded in popularity with [30] and led to fully convolutional encoder-decoder architectures being used in the medical domain, such as with U-Net [9], representing a major revolution in image segmentation and offering a leap in performance.

U-Net specifically constituted a breakthrough in medical imaging due to its ability to segment large images with high compute and data efficiency, and across multiple scales by introducing skip connections between encoder and decoder pathways. These allow the network to combine low-level spatial detail with high-level semantic context, enabling the segmentation of thin structures that require context such as cells and vasculature. The structure of U-Net can be extended into 3D with different approaches such as [31]. Recently, transformer-based architectures have superseded convolutional networks as a more generalist approach to machine learning, being applied to object detection [32] and being extended to segmentation, offering higher performance thanks to a more general computation model:

expert crafted feature extraction is removed, such as in the locality prior of convolutions, enabling the capture of more diverse features and long-range spatial dependencies, at the cost of an increase in the amount of training data required.

As algorithm complexity and data dimensionality increase, so must the amount of training or example data increase, an issue for convolutional networks and even more so for transformers. This limits the application of deep learning to Micro-CT data where annotated data is expensive and where inter dataset variance is large, even though attempts exist to palliate the data thirst by offering pre-trained models for 3D CT data [33] or by crafting more data efficient models. Examples also exist of trying to make use of so called foundational models, based on a large transformer and able to be “prompted” to customize the segmentation to the usecase at hand [34].

Deep learning has been applied to vascular segmentation in for blood vessel extraction directly with notable success: DeepVesselNet [12] introduced a family of architectures designed for vessel segmentation, centerline prediction, and bifurcation detection in 3D angiographic data by making explicit use of the structural priors inherent to blood vessels as secondary learning targets. Training on centerline and bifurcation labels jointly with segmentation masks improves topological results when compared to simple segmentation training, by providing extra prior generating signals to the model, highlighting the benefits of explicitly utilizing structural priors. V-Net [35] extended U-Net [9] into 3D and improved performance by modifying the loss function; Dice loss is used for the overlap between predicted and ground truth segmentation masks as opposed to cross-entropy, as it is better suited for the class-imbalanced setting characteristic of vessel segmentation where vessel voxels represent a small minority of the image volume.

Finally, hybrid approaches offer the interesting property of combining classical extraction methods with deep learning by leveraging input filters on the image to obtain richer features, such as applying vesselness maps to the image before processing, to integrate the priors of the vesselness filtering explicitly into the algorithm and reduce data needs [36].

Problem 3. *Segmentation evolved from classical algorithm driven methods towards data driven algorithms, without a corresponding increase in the availability of annotated CECT datasets. Our segmentation method should not make use of minimal external data sources to avoid requiring annotation, and make use of the prior work in algorithmic extraction.*

Unreliability of ground truth

Manual annotation of vascular networks in micro-CT images is time intensive and requires expert knowledge. When annotation is carried out, the resulting annotations are known to exhibit significant inter-annotator variability [8]. This variability represents an irreducible uncertainty in the ground truth that places an upper bound on the performance achievable by any supervised method, especially those training on voxel level data. Work-arounds exist, such as removing the inter-annotator noise by first performing structural extraction via skeletonization on the data, and then using these structures to segment the base image,

however this presents the problem of introducing the biases of the skeletonization step backwards in the pipeline. Annotations thus cannot be considered a definite truth, and as a result errors calculated from these annotations should take into account this potential variability.

Problem 4. *Manual ground truths can present an irreducible level of uncertainty, our method of evaluation should avoid reliance on pixel wise data.*

2.3.3 Data limitations

The issue of data availability for training and evaluating is well known, with a variety of approaches to tackle the issue: simulation can be used to generate data, using either geometrically driven techniques such as L-Systems, to generate branching tree structures grounded in the physiological laws of arterial branching, reviewed in [37], or neural-network driven systems such as generative adversarial networks (GANs) [38].

When data is available it may be used directly, although models trained on one imaging protocol, contrast agent, sample batch, or tissue type exhibit degraded performance when applied to data from a different distribution: this is termed domain shift, and is particularly acute in methods with imaging methods containing many adjustable parameters and machine specificities like micro-CT imaging. This can require the use of techniques like transfer learning, where a model is first trained on a similar task with large data availability then fine-tuned on a small amount of in-domain data [12]. Available data of any form may also be augmented: by degrading, deforming or otherwise modifying it without changing the annotations, robustness to certain shortcomings such as discontinuity can be improved at the potential cost of increased noise.

2.3.4 Error calculation

When training a model, or evaluating a method, it is important to be able to measure performance in a repeatable and objective way independent of human feedback. In their simplest form for binary classification, the error rate is simply the amount of correctly classified examples, however as dimensionality of outputs increased, so did the complexity of evaluation methods.

Loss functions such as cross-entropy used for binary classification calculate the loss on a pixel by pixel basis, based on the predicted distribution. It can be used for segmentation [9], however it is important to recall that our prediction task carries a heavy imbalance due to many more pixels being the background class. With cross-entropy loss, each prediction is independent, meaning for our biased distribution there is a prior of predicting background and producing fragmented or incomplete vessel predictions. There exist loss functions created that integrate the class imbalance explicitly, such as Dice [39] however these do not integrate the prior of connectivity or any spatial relatedness: each prediction is treated as independent.

For vasculature specifically, breaks in connectedness can be difficult to reconnect downstream, motivating the creation of custom loss function in [40] called clDice, that integrate the prior of connectedness, and build it into the loss function for predictions. In [41] clinically-relevant vascular features are encoded into the loss function. When evaluating a segmentation method consideration of the downstream analysis of its use to extract relevant features, such as tortuosity and branching ratio, is important to consider. As a result, and driven by the needs expressed by the team members, outputs will be evaluated based on point-wise loss of user annotated points, as well as on a manually annotated pixel level baseline. The challenging parameter of connectivity, difficult to capture in the loss, will be manually measured by selecting target areas and reported separately.

Problem 5. *Evaluation of a segmentation method often ignore downstream tasks. Our usecase should take downstream performance into account by being exportable into a shared format, as well as being evaluated on relevant challenging scenarios.*

2.3.5 Existing pipelines and tools

Several software tools and pipelines have been developed to support vascular segmentation and analysis workflows. ITKTubeTK [42] offers a library of algorithms for tubular structure segmentation and graph extraction built on the ITK framework. The Vascular Modeling Toolkit (VMTK) [16] integrates in 3DSlicer and provides a user friendly comprehensive suite of tools for vascular segmentation and vascular extraction: centerline extraction and surface reconstruction. VesselKnife [43] provides an integrated pipeline targeting vessel segmentation, skeletonization, and graph extraction from micro-CT data. For analysis, SKAN [44] provides a well documented Python library for the quantitative analysis of skeleton graphs extracted from binary segmentation masks, enabling computation of branch length distributions, tortuosity, and network connectivity. SimVascular [45] extends vascular extraction into simulation, enabling downstream modeling of blood flow within reconstructed vascular geometries.

The landscape of vascular segmentation tools reflects a broader conflict throughout bio-informatics, medical informatics and bioimager: general-purpose deep learning segmentation frameworks offer good performance on the datasets they are trained on, but require large annotated datasets and offer limited interpretability or controllability as well as are not baked into tools used by researchers. Classical model-based methods such as Frangi filtering are more transparent and adjustable, built into existing tools and easy to use, but require manual parameter tuning and struggle with complex vascular geometries.

Problem 6. *Existing pipelines focus on large vascularization. Our software should fit into*

the niche of small vessel segmentation, where tools are not readily available or tailored to the unique challenges of micro vasculature.

Chapter 3

Problem statement

The task at hand requires analyzing existing collected CECT data that presents challenges in the form of gradients, disconnected vessels and low contrast, by means of a user friendly piece of software. This software must be re-usable and fit into the wider ecosystem of open source tools for working on 3D data being implemented as an extension of 3D Slicer

The task at hand covers a multistep pipeline with three key stakeholders:

- future end users of the tool on different data
- the wider scientific community for whom replicability and extensibility matter, and
- computer science researchers, for whom algorithm development and performance is the goal.

From this, the following problem statement is obtained:

TODO revisit: How may a software pipeline for the segmentation and extraction of blood vessel networks from CECT images be designed, such that it is accessible and usable by domain scientists such as biologists, for whom software is not a primary tool (**Problem 1**), where the software is robust to the challenges of CECT: intensity gradients and variable contrast levels, without requiring dataset-specific hyper parameter tuning (**Problem 2, 4.2, 4.3**), capable of producing outputs in portable binary masks for use downstream (**Problem 3.1**), that is evaluated in a robust replicable fashion (**Problem 3.2, Problem 4.1**)?

3.1 Problem resolution approach

The following is proposed to approach the problem:

1. Use 3DSlicer [15], an open source, extendable software as a basis for a plugin
2. Placing the user in the segmentation loop and enforcing data driven hyperparameter selection through the placement of anchors [28]
3. The use of generalist, robust algorithms
4. Output in a portable binary segmentation format as individual slices or DICOM
5. The inclusion of user feedback during the development of the plugin

3. PROBLEM STATEMENT




Figure 3.1: Preliminary synthesized diagram for the annotation of vascularization and sharing of results

Chapter 4

Software tool development

The team had prior experience in three tools for the analysis of 3D data: Avizo: a proprietary and paid software, Dragonfly3D: with either a “FreeD” free-for-academics or paid commercial license, and a bare metal approach using python code. These three analysis methods were favored by different users: those having been in the team a long time had adapted to the team standard of Avizo, also the recommended starting point for new joiners; Dragonfly3D was used by team members that had previously worked on data elsewhere and used software on their own devices, and the “bare metal” approach was taken by students wishing to avoid the substantial learning curve and friction involved with using one of the two aforementioned softwares to carry out basic analysis. Avizo and Dragonfly3D were considered as potential candidates for the creation of a vessel extraction tool, however this proved challenging would not satisfy the desire for an open source alternatives offering extendability and re-usability by other researchers.

In order to maintain interoperability with the Avizo and Dragonfly3D the 3D Slicer plugin needed to export data in one of two ways: **(i)** a fashion that transparently fits into the existing pipeline, namely that of exporting raw binary segmentations that could then be re-imported elsewhere **(ii)** an export format that better enforces scientific rigor and enables easier data sharing, in the form of DICOM. This second option will offer the team growth perspectives in opening the door for easier collaboration with computer scientists as well as other researchers downstream.

4.1 3D Slicer Plugin

As the 3D Slicer ecosystem is diverse with many existing tools, available plugins with similar goals (vascular extraction) and similar data formats (high resolution CT) were tested. It was noted that many available plugins were either unmaintained, throwing errors during installation or use or did not successfully produce outputs; this is despite 3D Slicer having two classes of plugins: “official” via the extension manager [46] (further broken down by maintenance type) and user installable or “non official”.

After testing various plugins, the following common limitations came to surface:

- 3D Slicer hangs when executing algorithms, and this execution happens without communication. The OS will ask the user if they wish to stop the program.
- UI development is restrictive: UI must be in the left corner, and few context menus are available

- Execution in Python is inefficient, and plugins that aimed to use multiple cores called external libraries
- Machine learning extensions did not come pre-packaged with the model weights and often required a GPU and installation of external software
- Errors during installation were not user-friendly enough for a non computer scientist to be able to debug an installation problem
- The 3D Slicer forums contained ample amounts of users having issues with different plugins

4.1.1 Implementation

3D Slicer is multi platform and is built around plugins to carry out tasks, communicating with the Medical Reality Markup Language (MRML) Figure 4.1, where modules read and write to MRML nodes. These plugins can be implemented in one of three forms: as command line interface (CLI) tools where they may be called as an encapsulated piece of code and is the most abstracted way of working with external code or a C++/Scripted (Python) loadable module [1]. The Scripted loadable approach was selected for its use of Python, lowering the barrier to entry, ability to use the UI features of 3D Slicer, and having access to the full slicer API. Work was carried out in Visual Studio Code under Ubuntu 24.04.




Figure 4.1: Left: , Right: QT user interface designer

To implement a plugin, the extension wizard² was used to generate the required extension boilerplate, and the ScriptedLoadableModule was used as a starting point for the implementation code. 3D Slicer implements a user interface development module in QT, enabling near drag and drop creation of UI elements, making it the natural method of developing a seamless UI experience. Although basic, this radically simplifies the process of creating a cross platform UI.

²https://slicer.readthedocs.io/en/latest/user_guide/modules/extensionwizard.html#extension-wizard




Figure 4.2: Structure for module implementation from [1].

4.2 Segmentation of tumor vascularization

As seen during the literature review, a wide variety of solutions exist for segmentation, as well as a diversity of different plugins for 3D Slicer. Additionally, the problem of lacking existing annotated data limits the possible data driven approaches beyond those making use of basic annotations able to be done by the user. As a result, after testing existing plugins, it was chosen to implement a multi stage pipeline that would leverage user feedback.

4.2.1 Principled hyper parameter selection

During user interviews it was noted that software users often made decisions without a principled, replicable and data driven decision method. Windowing and grey value were selected based on expertise, problematic for multiple reasons:

1. As the selected values are not data driven, decisions and pipelines cannot easily be replicated, and are more fragile
2. As the user acts as the evaluator to the actions done, they induce bias in the subsequent steps
3. Any user without knowledge of the downstream steps or what to look out for will not be able to make a fully informed decision for the hyperparameters to select, requiring iteration and potentially falling into local optima.

A principled approach was required: in order to select thresholds and evaluate the performance of the algorithm “in the loop”, as well as offer the user immediate feedback, a system of anchor points was implemented. These anchors are the first step in the process, and are obtained by having the user click “add” followed by placing the point in any of the 2D

windows. They can be saved and exported, as well as imported. Following the learning that users often do not take the time to properly name files and folders, the name when exporting is automatically constructed from the current data series name and description, as well as the full current date and time.

A second key element implemented to guide downstream steps was the expected vessel size and deviation in voxels. This parameter is easy to measure and is critical for certain steps such as the subsequent seed annotation analysis.




Figure 4.3: View of the seed annotation and vessel size definition panes. The user may press add, click on the locations in any of the right hand panes, as well as import previous annotations or export the current points.

4.2.2 Thresholding segmentation

As noted in Chapter 2.3, the simplest method of segmentation is thresholding. Given the nature of CECT intends to give the structures of interest high grey values, this is an intuitive method with a strong prior. It however does not work in practice for a few reasons: **(i)** the shell effect in CECT, where the contrast enhancing agent has higher concentrations on the outside or surface of the tumor, results in values that would be segmented as “vessel” even though they are not. **(ii)** grey value gradients appear between the outside and center of the samples and **(iii)** incomplete staining resulting in discontinuous vessels.




Figure 4.4: Illustration of the shortcoming of threshold based segmentation, with the “shell” being included

The shell effect can be tackled by fitting a shape to the outside of the sample, however following a user interview demonstrating the feature a critical issue was raised: in tumors, it is common for the outside to contain large and plentiful vascularization. Removal would both bias results as well as reduce the overall performance by preventing these outside vessels from being segmented.

Secondly, a threshold does not hold up to the gradients in images that are present which, in prior work carried out on this data, resulted in rejection of samples with a gradient considered too large. Finally, standard thresholding does not do anything to combat the incomplete staining resulting in discontinuous vessels. As a result, manual thresholding as well as automated thresholding such as Otsu were rejected.

4.2.3 Traditional algorithms

When annotated by the user, vessels are identified on a local scale using the grey value difference between the vessel with contrast enhancing agent and the background. Grey value alone is, as laid out above, not *sufficient*, but it contributes some proof towards the presence of a blood vessel, given the presence of a grey value difference and a vessel-like shape. As a result, to obtain vessels from the user placed starting points, a series of different steps were assembled to form a *pipeline*.

To go beyond the limitations of threshold based segmentation, methods that make use of the blood vessel prior (algorithms designed with blood vessel or tubular structure extraction in mind) were tested: Frangi (or its generalization in SimpleITK) is a tubular prior algorithm widely available and high performance, being available as a C++ implementation. The algorithm is generally applied at multiple scales (different expected vessel sizes) with the outputs collected together into one segmentation. Frangi is known for having multiple hyperparameters that require fine tuning, in order to offer the user a simple solution, the parameters for the minimum and maximum vessel scale are pre-defined, and the vessel parameters hard coded.

Additionally, as the user has already provided high confidence starting points in the form of annotations, the marching squares algorithm was applied to extract vessels attached to these points, as mentioned in Lesage *et al* [22].

In order to combine the information provided by thresholding with that of Frangi, a novel approach was utilized: the use of a **probability map**.




Figure 4.5: Visualization of the vesselness probability map, as defined by the Frangi feature.

Higher valued pixels are assigned a higher vesselness probability.

This approach was selected as it allows the accumulation of evidence across steps in the pipeline, and if multiple probability maps are saved (one for each step or method) they can be given weights when being combined to create the final segmentation.

4.2.4 Machine learning

bla bla about the random forests

bla

bla

4.3 Results

After the first round of development, a multi stage pipeline was obtained: from the user placed points, a shell removal step was done, and subsequently the user “vessel” points

were grown using region growing to follow the ridges defined by the high grey value points. These points grown from the user annotations were then treated as a baseline for creating a simple classifier based on random forests: the vessel points, and points within a small region around them, were used. The probability maps generated from this step, alongside the tubular frangi, were then combined and thresholded based on the values of the user annotated points. This complex method led to promising results:




Figure 4.6: Visualization of the segmentation following the multi stage pipeline during development

The outputs from this stage presented interesting characteristics: continuous vessels and few noisy sections. Performance was measured using the user annotated points, with 26/27 vessels and 16/16 background points correctly classified.

Chapter 5

Contending with limitations

5.1 Micro-CT Data management challenges

Data management for Micro-CT scans is a challenge for users: after a scan is completed, they receive data from the CT machine in the form of a collection of 16 bit TIFF files: heavy, with a single 2000x2300 slice at 16bits per pixel weighing **9.2MB**, or as is often the case the data is saved as 3 channels, resulting in 27.6MB, and a whole 2400 slice scan weighing at least **22.1GB**. Scans are then windowed to 8 bit, occasionally with some form of compression, and the empty slices are removed: this generally halves or more the total data amount. This windowing process was documented as being unprincipled: the window was chosen based on the researchers best judgment, and the original uncompressed data discarded.

Furthermore, certain researchers would then carry out a lossy compression of the data in the form of JPEG image slices, as was the case with the data used in this thesis: the total scan weights provided ranged from **0.103** to **13.2GB** (597x698x854 to 3000x3000x2653) and the original lossless data was not preserved, in both cases the windowing and the compression were motivated by data storage cost concerns.

Finally, the provided data was generally given with little or no context: the data was provided in the form of a folder containing images as well as experiments that were run, with no associated dates and without grounding context such as the scan voxel size or parameters of the scanning machine. These issues of dataset size and compression resulted in challenges unforeseen during the literature review which required particular attention.

5.2 Performance challenges

The initial plugin development phase took place using the CA-LL-R dataset, the smallest uncompressed dataset Figure 5.1, to enable rapid prototyping and evaluation of the performance. Following the initial success in extraction based on points placed by the user, tests were carried out on larger datasets.




Figure 5.1: TODO: This is a placeholder. TODO: Add horizontal lines for the RAM of computers.

Visualization of the raw dataset sizes, obtained by multiplying width, height and depth by 8 bits per pixel

As noted in Chapter 4, the total data required for an uncompressed scan can reach into the tens or hundreds of GB. During the initial software evaluation, 3D Slicer was successful in loading all datasets on the development machine - however it was not verified at the time how much memory was being used. The testing of the pipeline on other datasets revealed the performance limitations of the implemented approach: with initial end to end runtime being about an hour and requiring 24GB of system memory, larger datasets saw an increase in inference time to un-manageable levels, as well as limitations of system memory.

These limitations performance issues have multiple sources: when implementing a 3D Slicer plugin in python, a single thread is available, and this thread locks all other 3D Slicer activity (this fact extends to other 3D Slicer functions such as loading and saving). When running on a large scan, combined with the generation of probability maps and the sequential algorithms, memory usage exceeded ram, reached into swap, and could run

seemingly indefinitely (success was only observed on smaller scans). This is a known issue with 3D Slicer³.




Figure 5.2:

1. CA-LL-L2 RAM utilization during segmentation: baseline after loading dataset, below 24GB utilization, during processing 100% RAM and swap are used, 2. Run 1, CB-LURU-R showing full RAM and SWAP utilization

Methods exist to handle such large datasets: the most basic approach is cutting down of full scans into smaller chunks, or subsampling the scans with some form of interpolation. Cutting scans down has the disadvantage of requiring stitching after running algorithms, and if done using 3D slicer's built in slicing, requires the ability to load the full dataset. An experiment was run, where a target scan was cut into 4 smaller sections using Python; this proved to be unwieldy for annotation and running the pipeline. Subsampling requires the the target structures to be large enough to allow it: subsampling to 1/4 resolution means that any vessel 4 voxels across would be reduced to approximately a single voxel.

Industry standard methods exist for handling large datasets: HDF5 is intended to *store* such large multidimensional arrays and efficiently enable loading subsections, however this data storage format is totally incompatible with most medical imaging software, and is not the standard used by CT machines. Standards such as DICOM did not provision for the possibility that data such as those generated by micro-CT may exist in the future in the medical domain, and do not deal well with dynamically loading large datasets from memory. To *operate on* large multidimensional arrays in Python, there exist libraries such as Dask

³Performance limitations as discussed on here the forums

that enable “chunking” of the data to process smaller areas: this could enable improved scaling to larger scans.

These performance concerns highlight a continuous issue encountered during the writing of this thesis: the complexity of methods able to be tested was limited by the choice of software, volume of data and the hardware available. Lab computers available to students have 32GB of ram, less than the computer used for the testing and writing of code, and it was noted by previous students working on MicroCT imaging that they had struggled to run algorithms across the whole image. In the end, much effort was invested in the research, testing and optimization of the algorithms, and runtime concerns pushed development towards the use of methods implemented in C++ available with Python bindings, such as the SimpleITK Frangi filter used.

5.3 Ablation study

Given the initial positive outputs of the pipeline but with the limitations in terms of memory and performance, it was decided to carry out an ablation study. It is common for software approaches comprising multiple steps to not correctly quantify the contribution of each individual step to the overall algorithm performance, and properly understand the performance/cost tradeoff of these steps.

In order to ablate the algorithm, the weights of the probability map of each step was set to 0 except the one under test, which was done for all the parameters of the pipeline. For performance measurements, RAM and CPU utilization were monitored alongside runtime with unused steps disabled. This showed that the Frangi vesselness step was the most important to obtaining high vessel extraction performance, offering the same 26/27 correctly classified vessel points as the full pipeline. Additionally the performance evaluation method was highlighted as lacking: false positives and discontinuities in the blood vessels were not properly penalized, and small vessels were not being correctly segmented; all things that point wise annotations fail to capture.

5.4 Conclusion of the ablation study

The removal of all steps (shell removal, random forest) except vesselness resulted in improved performance for extraction of vessels in the outer shell, as noted previously being a crucial point, as well as a large performance increase: vessel extraction went from taking multiple hours to under 20 minutes. It also successfully reduced RAM and swap usage, enabling effective running of the algorithm on larger samples. As a result of this small ablation study, the probability map approach was simplified to reduce memory usage, with a single map being successively updated, and focus was moved towards the use of simple, scientifically grounded extracton based on line-like features with a step for combatting the disconnections in vessels.

Method	CPU Threads	RAM GB	RAM Peak GB	Performance Contribution	Inference Time
Baseline	1	9	NA	NA	
Loaded dataset & point annotation	1	13	NA	NA	
Pre-processing (shell, intensity)	1	13.4	19	24 of 27	19
Vesselness	20	22.5	38.4	26 of 27	17
Fast marching algorithm	1	28	28	Minor	22
Tube extraction	1	20	21	Minor	1786
Tree training and inference	1	16	32	Minor	1991

Table 5.1: Ablation study measurements

Chapter 6

Micro vasculature extraction

6.1 Final vessel extraction flow

6.2 Data annotation for evaluation

As mentioned, the user is expected to place vessel, background and outside-of-volume points. These act as points used to define hyperparameters of the algorithms, but also as a performance metric: when the pipeline is run, feedback is given with how many vessel points are correctly classified. However this method of performance evaluation has shortcomings: it evaluates the data in a pointwise fashion, ignoring critical elements for downstream tasks such as connectivity, and relies on the human evaluating a 2D plane, ignoring parameters such as gap filling. As users also place points generally towards the center of the vessels, there is little measurement of the width of vessels beyond if a background point ends up being caught in the vessel prediction.

As a result of this, it was decided to provide more dense annotations in the form of fully annotated regions taken from different scans:

6.2.1 Annotation procedure & tools

Annotation was carried out using 3D slicer on 64x64x64 regions randomly selected from the provided samples: these regions were selected from one of 25 regions cropped from the center of the image. This ensures that points exist

6.3 Performance results

The final algorithm was ran on 192x192x192 volumes to provide context around the annotated 64x64x64 regions:

Chapter 7

Conclusion and future perspectives

discussion about the results, reposition the work in the context, plan future steps

I conclude that all is well that ends well

Appendices

Appendix A

Appendices

1.1 Environmental and CO2 impact of the thesis

The main sources of CO2 impact of this thesis are the use of computational resources and the human factor.

1. During the thesis, the laptop of the writer was used: A legion 7i slim with 48GB DDR5, i7-13700H, 8GB RTX4070.
 1. **Production CO2:** The compliance document provided by Lenovo details an estimated carbon footprint of 429g +/- 86g CO2e. For a lifetime of 8 years, and an estimated thesis length of 3 months equivalent full time, this equates to **13.41Kg** CO2e.
 2. **Usage CO2:** The total estimated time dedicated to the coding and writing of the thesis is estimated based on the credits and approximate time investment per credit: 25 credits at 30 hours per credit equates to 750 hours. The laptop was measured over one 2h coding + writing session as using 40.89Wh, equating to a continuous use of 20.45w. This works out to 15.34 kWh, with an estimated Belgian CO2/kWh by the AIB (Association of Issuing Bodies) in 2024 of 131.73 g/kWh, resulting in total emissions of **2.02Kg** CO2.
 3. **AI Utilization:** The impact of LLMs from large providers such as OpenAI is an ongoing research topic and difficult to measure, on top of which are layered issues like the use of Piccolo, UCLouvains aggregation service. A total of TODO credits were used on Piccolo, with an unknown impact. Claude Code was also used, with a total of 3 separate conversations containing TODO 41 cumulative prompts. The impact of AI use is thus left out of the estimation, with a guesstimate below.
2. The thesis required meetings, for which the main source of CO2 emissions is the travel to/from Louvain-La-Neuve. This travel was done primarily by train, and the amount of travel directly attributable to the thesis was of 21 round trips, with a distance per trip of 93.2Km at an estimated 16.6g CO2e/Km, equating to **32.49Kg** CO2e.

There was minimal extra data storage required to complete this thesis, however it is important to note that the large files handled do incur emissions if duplicated and stored across devices outside the user's own computer.

The prior calculations equate to a total approximate impact of **47.92 Kg CO2e**.

As for AI - a rough impact is estimated using the computer of the writer as a basis: three of the prompts sent to Piccolo and Claude were sent to the largest local qwen instance fitting in the 8GB GPU, and the time to answer as well as power use while answering estimated: for the three prompts, a total of TODO minutes to answer was required, with a power use of TODO

1.2 Instructions for scientific rigor in data handling

1.3 Use of LLMs and AIs

The main two AI assistants used in this thesis were Piccolo, offered by UCLouvain and used for writing feedback, and ClaudeCode, by Anthropic, for the expansion and creation of the plugin.

1.4 Structuring a Masters thesis

During the preparation, reading and writing phases of the masters thesis, it was particularly challenging to both understand what is expected as well as the techniques needed for writing scientifically. This process of learning to write a scientific article has been challenging as well as rewarding, but required extensive effort to document and learn beyond what was offered by the official UCLouvain Masters thesis page.

The writer strongly recommends the *Harvard College Writing Center's* recommendations for thesis writing, as well as the talk *How to write your PhD thesis (without going insane)* by James Hayton.

1.5 Masters thesis writing across group and field boundaries

As is all too common, it is easy when experience is lacking to over estimate ones capabilities, as well as the ease of solving (or even working on) a particular problem. The concept of a thesis across field boundaries was initiated by the writer motivated by the wish to grow a skillset beyond what is usually developped when carrying out a thesis within ones field, as well as by a desire for real world impact through applying computer science skills to a problem where the limitations are mainly technical.

This theis was particularly challenging as a result, not due to educational or technical reasons, but mainly due to the practical issues one faces when attempting to solve a real problem: unclear or varying goals and measurements for success, a misunderstanding by the engineer of the field in which the software is to be used, as well as the friction induced by having many stakeholders. These three points will be unpacked one by one:

1.5.1 Setting clear goals

1.5.2 Field knowledge and problem encapsulation

1.5.3 Stakeholders

Appendix B

Extra figures

Bibliography

- [1] L. for Percutaneous Surgery, “Tutorial on 3D Slicer python scripting and programming.” [Online]. Available: <https://github.com/Slicer/SlicerProgrammingTutorial>
- [2] T. Balcaen *et al.*, “Revealing the Three-Dimensional Murine Brain Microstructure by Contrast-Enhanced Computed Tomography,” *Frontiers in Neuroscience*, vol. 17, p. 1141615, 2023, doi: [10.3389/fnins.2023.1141615](https://doi.org/10.3389/fnins.2023.1141615).
- [3] S. de Bournonville, S. Vangrunderbeeck, and G. Kerckhofs, “Contrast-Enhanced MicroCT for Virtual 3D Anatomical Pathology of Biological Tissues: A Literature Review,” *Contrast Media and Molecular Imaging*, vol. 2019, p. 8617406, 2019, doi: [10.1155/2019/8617406](https://doi.org/10.1155/2019/8617406).
- [4] G. S. G. S. J. H. N. C. C. B. S. E. H. B. M. Anna K. Hans D., “Extending two-dimensional histology into the third dimension through conventional micro computed tomography,” *NeuroImage*, vol. 139, pp. 26–36, 2016, doi: [http://dx.doi.org/10.1016/j.neuroimage.2016.06.005](https://doi.org/10.1016/j.neuroimage.2016.06.005).
- [5] J. Pichat, J. E. Iglesias, T. Yousry, S. Ourselin, and M. Modat, “A Survey of Methods for 3D Histology Reconstruction,” *Medical Image Analysis*, vol. 46, pp. 73–105, 2018, doi: [10.1016/j.media.2018.02.004](https://doi.org/10.1016/j.media.2018.02.004).
- [6] A. Maes *et al.*, “Cryogenic contrast-enhanced microCT enables nondestructive 3D quantitative histopathology of soft biological tissues,” *Nature Communications*, vol. 13, no. 1, p. 6207, Oct. 2022, doi: [10.1038/s41467-022-34048-4](https://doi.org/10.1038/s41467-022-34048-4).
- [7] L. Wlodarski, “Contrast-enhanced micro-CT for 3D X-ray based histology to assess the influence of antiangiogenic drugs on tumor vascularization and necrosis.” 2024.
- [8] D. Lin *et al.*, “\textit{E pluribus unum}: prospective acceptability benchmarking from the Contouring Collaborative for Consensus in Radiation Oncology crowdsourced initiative for multiobserver segmentation,” *Journal of Medical Imaging*, vol. 10, no. S1, p. S11903, Feb. 2023, doi: [10.1117/1.JMI.10.S1.S11903](https://doi.org/10.1117/1.JMI.10.S1.S11903).
- [9] “U-Net: Convolutional Networks for Biomedical Image Segmentation,” 2015, doi: <https://doi.org/10.48550/arXiv.1505.04597>.
- [10] C. T. C. Inc., “Dragonfly FreeD License Model.” [Online]. Available: <https://dragonfly.comet.tech/en/products/license-models/freed-license>
- [11] S. Jodogne, “The Orthanc Ecosystem for Medical Imaging,” *Journal of Digital Imaging*, vol. 31, no. 3, pp. 341–352, 2018, doi: [10.1007/s10278-018-0082-y](https://doi.org/10.1007/s10278-018-0082-y).
- [12] G. Tetteh *et al.*, “DeepVesselNet: Vessel Segmentation, Centerline Prediction, and Bifurcation Detection in 3-D Angiographic Volumes,” *Frontiers in Neuroscience*, vol. 14, p. 592352, 2020, doi: [10.3389/fnins.2020.592352](https://doi.org/10.3389/fnins.2020.592352).

- [13] Y. He *et al.*, “SPROUT: A User-friendly, Scalable Toolkit for Multi-class Segmentation of Volumetric Images,” *bioRxiv*, 2024, doi: [10.1101/2024.11.22.624847](https://doi.org/10.1101/2024.11.22.624847).
- [14] I. O. team, “Frangi Vesselness – ImageJ.”
- [15] K.-C. J. F. J. F.-R. J.-C. P. S. B. C. J. D. F. F. S. M. B. J. A. S. M. J. P. S. K. R. Fedorov A. Beichel R., “3D Slicer as an Image Computing Platform for the Quantitative Imaging Network,” *Magnetic Resonance Imaging*, vol. 30, no. (9), pp. 1323–1341, Nov. 2012, doi: [10.1016/j.mri.2012.05.001](https://doi.org/10.1016/j.mri.2012.05.001).
- [16] R. Izzo, D. Steinman, S. Manini, and L. Antiga, “The Vascular Modeling Toolkit: A Python Library for the Analysis of Tubular Structures in Medical Images,” vol. 3, no. 25. p. 745, 2018, doi: [10.21105/joss.00745](https://doi.org/10.21105/joss.00745).
- [17] A. Affane *et al.*, “The R-Vessel-X Project,” *arXiv preprint*, 2025.
- [18] L. He, L. R. Long, S. Antani, and G. R. Thoma, “Histology Image Analysis for Carcinoma Detection and Grading,” *Computer Methods and Programs in Biomedicine*, vol. 107, no. 3, pp. 538–556, Sept. 2012, doi: [10.1016/j.cmpb.2011.12.007](https://doi.org/10.1016/j.cmpb.2011.12.007).
- [19] P. B. N. M. B. C. M. A. K. & F. P. Juliana Martins de S. e Silva Irene Zanette, “Three-dimensional non-destructive soft-tissue visualization with X-ray staining micro-tomography,” 2015.
- [20] Z.-M. S. D. C. Ko-Chin Chen Alon Arad, “High-definition neural visualization of rodent brain using micro-CT scanning and non-local-means processing,” *BMC Med. Imaging*, vol. 18, p. 38, Oct. 2018, doi: [10.1186/s12880-018-0280-6](https://doi.org/10.1186/s12880-018-0280-6).
- [21] G. K. Lisa Leyssens Camille Pestiaux, “A Review of Ex Vivo X-ray Microfocus Computed Tomography-Based Characterization of the Cardiovascular System,” 2021.
- [22] D. Lesage, E. D. Angelini, I. Bloch, and G. Funka-Lea, “A review of 3D vessel lumen segmentation techniques: Models, features and extraction schemes,” *Medical Image Analysis*, vol. 13, no. 6, pp. 819–845, 2009, doi: <https://doi.org/10.1016/j.media.2009.07.011>.
- [23] F. G. S. B. D. M. J. M. S.-A. R. H. G. M. N. R. K. W. L. F. K. T. L. Josef Ehling Benjamin Theek, “Micro-CT Imaging of Tumor Angiogenesis Quantitative Measures Describing Micromorphology and Vascularization,” *Am. J. Pathol.*, vol. 184, pp. 431–441, Feb. 2014, doi: [10.1016/j.ajpath.2013.10.014](https://doi.org/10.1016/j.ajpath.2013.10.014).
- [24] “A Threshold Selection Method from Gray-Level Histograms,” 1979, doi: <https://doi.org/10.1109/TSMC.1979.4310076>.
- [25] M. Mancas, B. Gosselin, and B. Macq, “Segmentation using a region-growing thresholding,” in *Image Processing: Algorithms and Systems IV*, E. R. Dougherty, J. T. Astola, and K. O. Egiazarian, Eds., SPIE, 2005, pp. 388–398. doi: [10.1117/12.587995](https://doi.org/10.1117/12.587995).
- [26] “Multiscale vessel enhancement filtering,” 2012, doi: [10.1145/3065386](https://doi.org/10.1145/3065386).
- [27] “Beyond Frangi: An improved multiscale vesselness filter,” 2023, doi: [10.1117/12.2081147](https://doi.org/10.1117/12.2081147).
- [28] “Tubular Structure Segmentation Based on Minimal Path Method and Anisotropic Enhancement,” 2011, doi: [10.1007/s11263-010-0331-0](https://doi.org/10.1007/s11263-010-0331-0).
- [29] T. Chen and C. Guestrin, “XGBoost: A Scalable Tree Boosting System,” *CoRR*, 2016, [Online]. Available: <http://arxiv.org/abs/1603.02754>
- [30] “ImageNet Classification with Deep Convolutional Neural Networks,” 2012, doi: [10.1145/3065386](https://doi.org/10.1145/3065386).
- [31] “3D U-Net: Learning Dense Volumetric Segmentation from Sparse Annotation,” 2016, doi: [10.48550/arXiv.1606.06650](https://doi.org/10.48550/arXiv.1606.06650).

-
- [32] N. Carion, F. Massa, G. Synnaeve, N. Usunier, A. Kirillov, and S. Zagoruyko, “End-to-End Object Detection with Transformers,” *CoRR*, 2020, [Online]. Available: <https://arxiv.org/abs/2005.12872>
- [33] S. A. A. K. J. P. & K. H. M.-H. Fabian Isensee Paul F. Jaeger, “nnU-Net: a self-configuring method for deep learning-based biomedical image segmentation,” *Nature Methods*, vol. 18, pp. 203–211, 2021, doi: [10.1038/s41592-020-01008-z](https://doi.org/10.1038/s41592-020-01008-z).
- [34] F. Semeraro, A. M. Quintart, S. F. Izquierdo, and J. C. Ferguson, “TomoSAM: A 3D Slicer extension using SAM for tomography segmentation,” *SoftwareX*, vol. 31, p. 102218, 2025, doi: <https://doi.org/10.1016/j.softx.2025.102218>.
- [35] “V-Net: Fully Convolutional Neural Networks for Volumetric Medical Image Segmentation,” 2016, doi: [??](https://doi.org/10.1016/j.media.2016.101552).
- [36] Y.-C. Chen *et al.*, “Coronary Artery Segmentation in Cardiac CT Angiography Using 3D Multi-Channel U-net.” [Online]. Available: <https://arxiv.org/abs/1907.12246>
- [37] M. Zamir, “Arterial branching within the confines of fractal L-system formalism,” *The Journal of General Physiology*, vol. 118, no. 3, pp. 267–276, Sept. 2001, doi: [10.1085/jgp.118.3.267](https://doi.org/10.1085/jgp.118.3.267).
- [38] X. Yi, E. Walia, and P. Babyn, “Generative adversarial network in medical imaging: A review,” *Medical Image Analysis*, vol. 58, p. 101552, 2019, doi: <https://doi.org/10.1016/j.media.2019.101552>.
- [39] C. H. Sudre, W. Li, T. Vercauteren, S. Ourselin, and M. J. Cardoso, “Generalised Dice overlap as a deep learning loss function for highly unbalanced segmentations,” *CoRR*, 2017, [Online]. Available: <http://arxiv.org/abs/1707.03237>
- [40] S. Shit *et al.*, “clDice - a Topology-Preserving Loss Function for Tubular Structure Segmentation,” *CoRR*, 2020, [Online]. Available: <https://arxiv.org/abs/2003.07311>
- [41] Y. Zhou *et al.*, “CF-Loss: Clinically-relevant feature optimised loss function for retinal multi-class vessel segmentation and vascular feature measurement,” *Medical Image Analysis*, vol. 93, p. 103098, 2024, doi: <https://doi.org/10.1016/j.media.2024.103098>.
- [42] “Initialization, noise, singularities, and scale in height ridge traversal for tubular object centerline extraction,” 2002, doi: [10.1109/42.993126](https://doi.org/10.1109/42.993126).
- [43] “VesselKnife – software for the analysis of tubular structures in biomedical images,” 2025, doi: [10.1016/j.cmpb.2025.108817](https://doi.org/10.1016/j.cmpb.2025.108817).
- [44] “A new Python library to analyse skeleton images confirms malaria parasite remodelling of the red blood cell membrane skeleton,” 2018, doi: [10.7717/peerj.4312](https://doi.org/10.7717/peerj.4312).
- [45] “SimVascular: An Open Source Pipeline for Cardiovascular Simulation,” 2016, doi: [10.1007/s10439-016-1762-8](https://doi.org/10.1007/s10439-016-1762-8).
- [46] 3. D Slicer Community, “Extensions — 3D Slicer documentation.” [Online]. Available: https://slicer.readthedocs.io/en/latest/user_guide/extensions.html#extensions-manager